

Facilitator guide

90 minutes | October 21, 2026 | 3:30-5:00 PM CT

Time	Slides	Focus / product
3:30-3:38 (8 min)	1-4	Frame the decisions. Identify one review task and rate a starting capacity.
3:38-3:50 (12 min)	5-7	Build a structured prompt. Draft a constrained prompt with an evidence requirement.
3:50-4:05 (15 min)	8-11	Demonstrate administrative review. Test synthetic application triage; catch omissions and unsupported claims.
4:05-4:20 (15 min)	12-14	Calibrate a review rubric. Compare independent human ratings before seeing an AI draft.
4:20-4:30 (10 min)	15-16	Protect the workflow. Place a data gate and a named human decision owner in a workflow.
4:30-4:45 (15 min)	17-20	Evaluate applicant AI use. Interpret a detector signal; draft a fair policy response.
4:45-4:55 (10 min)	21-23	Choose a capacity target. Select one capacity, evidence of progress, and a 30-day action.
4:55-5:00 (5 min)	24-25	Commit and close. Share an exit ticket and address remaining questions.

Before the session

Open the hub, practice page, and capacity matrix. Download the PPTX, PDF, workbook, and reference outputs. Test an approved Gen AI account with the synthetic packet only. Verify projection and Wi-Fi; print the workbook and matrix for participants without devices. Recheck detection guidance shortly before the workshop. Slide 26 is a source appendix, outside the 90-minute flow.

Demonstrations and debriefs

Keep both demos bounded; never substitute real applicant records.

3:38-3:50 | Prompt design

Slides 5-6: explain the five prompt components and compare the vague and bounded requests. Slide 7: give pairs 2 minutes to choose a task, 2 to revise, and 1 to critique. Ask which assumption they removed.

3:50-4:05 | Administrative triage

3:50-3:53: introduce P1-P5. 3:53-3:58: copy the triage prompt and synthetic packet from the practice page into an approved tool. Inspect missing enrollment confirmation, evidence IDs, and P5 instruction handling. 3:58-4:01: compare against the prepared reference. 4:01-4:05: pairs trace, challenge, and correct one output. If the tool fails, use the authored reference table; label it as such.

4:05-4:20 | Rubric calibration

4:05-4:09: explain the 0-2 anchors and insufficient-evidence distinction. 4:09-4:15: allow 2 minutes for independent ratings, 2 for comparison, and 2 for anchor revision. 4:15-4:20: run the rubric prompt on the same synthetic packet. Compare the AI draft only after human ratings. Expected discussion: goal and contribution support 2; reflection can expose differences in how reviewers interpret the anchor. Do not present a total or rank.

Debrief questions

What fact was unsupported? Which source location supports the claim? Did the output treat P5 as an instruction? Was disagreement about evidence or about the rubric? Who can stop the process? Keep a visible record of corrections rather than celebrating only fast generation.

Policy, capacity, and closure

Protect time for participant decisions.

4:20-4:30 | Protected workflow

Use 4 minutes for the data gate and 6 for mapping checkpoints. Ask pairs to name the owner before the tool, before reviewer use, and before a final decision. Require one explicit stop condition.

4:30-4:45 | Applicant AI use

3 minutes: explain detector limitations and study context. 3 minutes: use the hypothetical 1,000 x 1% = 10 example; this is not a product benchmark. 3 minutes: separate permitted use, disclosure, and misrepresentation. 6 minutes: groups analyze translation-support scenario, draft a neutral message, and name a review path. Debrief against the published rule, not impressions of writing style.

4:45-4:55 | Capacity target

3 minutes: revisit the six capacities. 5 minutes: choose a row, current evidence, and next action. 2 minutes: define a 30-day pilot with owner, measure, and stop rule. Ask participants to export their local plan or mark the printed matrix.

4:55-5:00 | Exit and questions

3 minutes: one prompt change, one safeguard, one policy decision. 2 minutes: questions and hub reminder. If running late, shorten whole-group reporting and optional explanation, not the independent rubric pass or the policy-response activity.

What to collect or retain

Participants keep their prompt, evidence corrections, rubric discussion, workflow map, policy starter, and capacity plan. No applicant information is collected by this workshop site. Do not collect participant browser exports unless there is an explicit agreed purpose.

Sources: NIST Gen AI Profile (2024); Turnitin detection FAQ (reviewed September 24, 2026); Liang et al., Patterns (2023); Al Ali, Helcl & Libovicky, EACL SRW (2026). See the hub for scope and interpretation.