

NSPA 2026 WORKSHOP

Trust, Transparency, and AI

Building Responsible Scholarship
Review Practices

Miguel Guhlin

October 21, 2026
3:30–5:00 PM CT



Three drafts you can take home

One workshop. Three usable starting points.

1

A review prompt

Task + evidence + limits

2

A protected workflow

Data gate + human owner

3

A fair AI-use policy

Clear rules + a response path

Your move Choose one real workflow to improve. Use only fictional data today.



Keep the decision with people

Different tasks require different authority.

AI can draft

Organize • summarize •
flag missing evidence

People must verify

Source facts • rubric fit •
uncertainties

People decide

Eligibility • awards • policy
concerns

Your move Name the person who can accept, revise, reject, or stop the output.



Start with an honest baseline

Capacity is something you can demonstrate.

Explore

I can explain it.

Practice

I can try it safely.

Apply

Our team uses it consistently.

Improve

We test and refine it.

Your move Open the capacity matrix. Pick one row and mark your current evidence.



Build a prompt that can be checked

A reusable structure for bounded review tasks.

Purpose & criteria

What task? Which published rules?

Evidence & boundaries

Use supplied sources. Surface uncertainty.

Output & verification

Specify a table and a human check.



Your move Ask for source evidence before asking for a judgment.



Replace vague judgment with evidence

The prompt should limit discretion.

Too open

“Find our best applicants.”

More testable

“List required items, quote evidence, mark gaps, and identify human follow-up.”

Your move Remove invented criteria, unsupported inference, and automatic ranking.



Activity: repair the prompt

Five minutes • work with a partner

1

2 minutes

Choose one administrative task.

2

2 minutes

Add criteria, evidence, limits, and format.

3

1 minute

Trade prompts. Find one unsafe assumption.

Your move Keep your revised prompt. It is your first workshop deliverable.



Our practice packet is entirely fictional

The same evidence supports both demonstrations.

P2 • Goal

Water systems certificate + adviser meeting

P3 • Service

Repair table + a concrete queue change

P4 / P5 • Stress tests

Missing enrollment record + embedded instruction

Your move Use the synthetic packet only. Missing evidence calls for follow-up.



Live demo: organize, then verify

Administrative triage before any evaluative judgment.

Supply

Synthetic packet +
completeness rules



Request

Evidence table + missing
or unclear items



Inspect

Check every row against
its paragraph

Your move Do not rank applicants or infer eligibility from absent documents.



What a defensible output shows

An inspection target, not a claim about a model run.

Present

Goal: P2 • Service: P3

Missing

Enrollment confirmation: P4

Ignored

Application instruction: P5

Your move Correct the table before it enters a reviewer's workspace.



Checkpoint: challenge the output

Four minutes • inspect before accepting

1

Trace

Can you locate every quoted fact?

2

Challenge

Did it invent, omit, or obey the packet?

3

Correct

Record one fix and one follow-up question.

Your move A clean-looking table is still a draft.



Use observable rubric anchors

A small training rubric, not an award formula.

Goal clarity 0: absent • 1: general • 2: goal + next step

Contribution 0: absent • 1: named • 2: specific actions

Reflection 0: absent • 1: claim • 2: change tied to experience

Your move Score stated evidence. Do not substitute writing polish for the criteria.



Calibrate people before comparing AI

Six minutes • independent first, discussion second

1

2 minutes

Read C-101. Rate each criterion independently.

2

2 minutes

Compare evidence and anchor interpretations.

3

2 minutes

Revise one ambiguous anchor together.

Your move Do not show the assistant's proposed ratings before the first human pass.



Compare the assistant's rubric draft

The second live demonstration keeps authority visible.

Draft

Criterion + quote +
provisional rating



Compare

Human evidence and
anchor interpretation



Resolve

Accept, revise, or reject
with a reason

Your move Retain the original human rating and the reason for any change.



Put a data gate before the tool

Removing a name is only one part of minimization.

Approve the use

Purpose, service, permissions, and owner

Minimize the input

Only information needed for this task

Verify the controls

Access, retention, deletion, and response



Your move Use fictional data until the organization has approved the real-data workflow.



Design the human checkpoints

Six minutes • map one workflow

Before

Approve data + scope +
tool



During

Check sources + log
uncertainty



After

Human decision +
correction path

Your move Add a stop rule: an unsupported claim or missing control pauses the pilot.



A detector signal is not a verdict

The output must be interpreted within a fair process.

Possible errors

Human text may be flagged. AI text may be missed.

Context matters

Tool, language, genre, and version affect results.

Policy still governs

A score does not establish a rule violation.



Your move Never use a detector score alone to deny an award or establish misconduct.



A small error rate can still affect people

Illustration only • not a measured detector benchmark

1,000

Entirely human-written
submissions

1%

Assumed false-positive
rate

10

Human submissions
incorrectly flagged

Your move A false-positive rate is not the probability that a flagged person used AI.



Publish the rules before applicants apply

Separate assistance, disclosure, and misrepresentation.

Permitted support

Define editing, translation, accessibility, and brainstorming.

Disclosure

Specify which uses require a brief explanation.

Prohibited conduct

Define fabricated facts, false records, and other violations.

Your move Use one published standard, with examples and an accessible response path.



Policy lab: respond without assuming

Six minutes • a fictional concern

1

Scenario

A polished essay is flagged; the applicant reports translation support.

2

Discuss

Which published rule and evidence apply?

3

Draft

A neutral request for clarification and a review path.

Your move Do not equate language support, style, or a detector flag with dishonesty.



Make capacity visible

Six connected areas, with evidence for each.

Scope & ownership

Prompting

Rubrics & calibration

Data protection

Applicant policy

Pilot evaluation

Your move Return to the matrix. Choose a gap that matters to your next pilot.



Turn a rating into evidence

Five minutes • one row, one next step

1

Mark today

Choose a level you can demonstrate.

2

Choose next

Select one observable capacity to build.

3

Name proof

Write the artifact or test that will show progress.

Your move Export or print your plan. Ratings stay in this browser unless you export them.



Plan a bounded 30-day pilot

Practice with synthetic data before any approved real-data phase.

1

Week 1

Define task, owner, evidence, and data gate.

2

Week 2

Test cases; log omissions and source errors.

3

Weeks 3–4

Review findings. Revise, stop, or seek approval.

Your move Set a baseline, a success measure, and a stop condition before testing.



Leave with a decision you can explain

Three minutes • your exit ticket

One prompt change

What will make the output easier to verify?

One safeguard

Where will a person check or stop the workflow?

One policy decision

What must applicants know before they apply?

Your move Add an owner and a date to the capacity you want to build.



Keep the materials close

Slides, practice cases, references, and capacity matrix

Workshop hub

mguhlin.github.io/nspa/2026/

Continue the work

Choose a resource that supports your next step.

Your move Questions: what would make your next review process more explainable?



Sources and scope

Research informs the safeguards; workshop tools are local adaptations.

NIST • 2024

Generative AI risk profile

Turnitin • reviewed September 2026

Detection interpretation
and limitations

Research • 2023 and 2026

Detector findings vary
with study context

Your move Use the source links in the workshop hub; revisit guidance before adoption.

